

```
In [49]: # Data Manipulation and Analysis
import pandas as pd
import numpy as np

# Visualization Libraries
import matplotlib.pyplot as plt
import seaborn as sns

# File Handling
from glob import glob

# Audio Processing
import librosa
import librosa.display

# Interactive Audio Playback
import IPython.display as ipd

# Utility Functions
from itertools import cycle

# Set Plotting Theme
sns.set_theme(style="white", palette=None)

# Define Color Palette and Cycle for Visualizations
color_pal = plt.rcParams["axes.prop_cycle"].by_key()["color"]
color_cycle = cycle(color_pal)

# SciKit Libraries (this is what your machine learning thing is derived from)
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report

In [50]: audio_file1 = glob("*/Users/Parag/Documents/THRT_data/noteA.mp3")
audio_file2 = glob("*/Users/Parag/Documents/THRT_data/noteB.mp3")
```

```
In [51]: #Audio_File1
print("Play the A note on the tuning fork at 440Hz")
ipd.Audio(audio_file1[0])
```

```
Play the A note on the tuning fork at 440Hz
```

```
Out[51]: 
```

```
In [52]: #Audio_File2
print("Play the B note on the tuning fork at 440Hz")
ipd.Audio(audio_file2[0])
```

```
Play the B note on the tuning fork at 440Hz
```

```
Out[52]: 
```

```
In [53]: #next Loading the audio files data
#first for the first audio file
Atime = Asr = librosa.load(audio_file1[0])
#displaying the first ten samples
print(f"The first 10 samples of note A (Atime:[10])")
print(f"The shape of Atime (Atime.shape)")
print(f"The sample rate for the A note file (Asr)")

The first 10 samples of note A [-1.1165938e-23  9.0989867e-24 -4.4460958e-24  1.9438744e-23
-6.6528909e-25 -1.6130022e-23 -6.6174649e-24 -6.4106497e-24
 2.4786012e-24  2.8237732e-23]
The shape of Atime (4433904,)
The sample rate for the A note file 22050
```

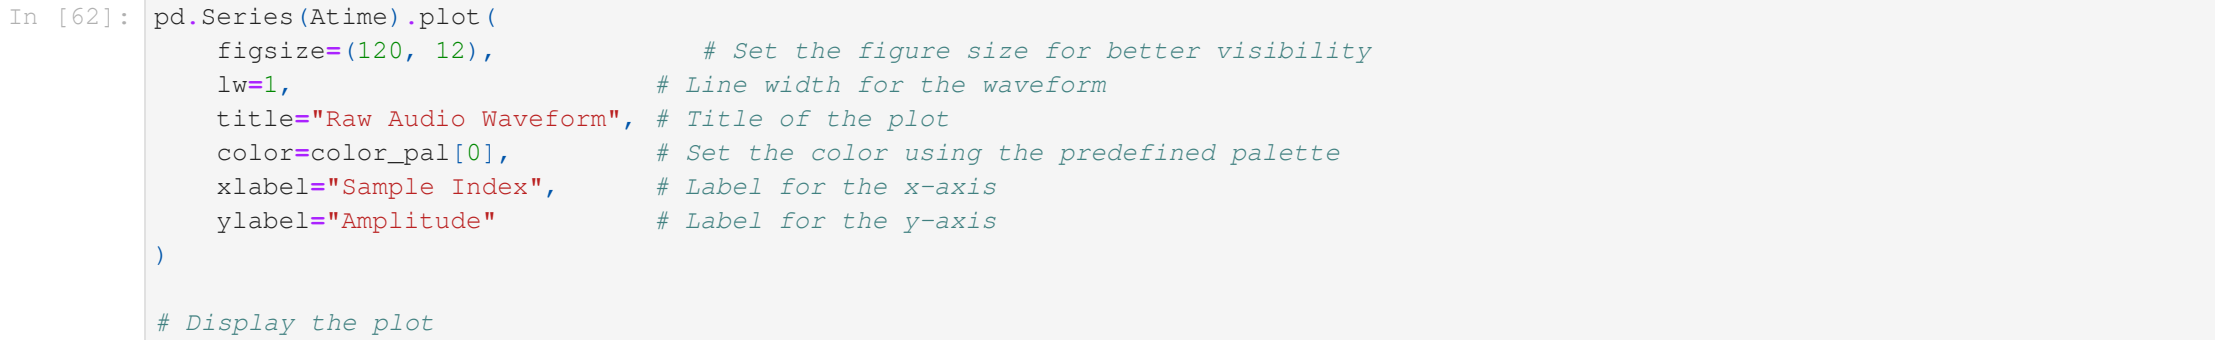
```
In [54]: #for the second audio file
Btime,Br = librosa.load(audio_file2[0])
#displaying the first ten samples,shape and sample rate
print(f"The first 10 samples of note B(Btime:[10])")
print(f"The shape of Btime (Btime.shape)")
print(f"The sample rate for the B note file (Br) in")
print(f"(*I looked up why i'm getting a bunch of zeros in the first then samples -> There's this thing called a DC offset which gets added)
```

```
The first 10 samples of note B[0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
The shape of Btime (4433904,)
The sample rate for the B note file 22050
```

```
I looked up why i'm getting a bunch of zeros in the first then samples -> There's this thing called a DC offset which gets added to a file beginning, librosa counteracts by subtracting the audio file with the mean of the samples and so we get a value that is very close to zero which gets approximated to zero
```

```
In [55]: #Plotting the raw audio waveform for the A note
# Plot the raw audio waveform
pd.Series(Atime).plot(
    figsize=(12, 7), # Set the figure size for better visibility
    lw=1, # Line width for the waveform
    title="Raw Audio Waveform", # Title of the plot
    color=color_pal[0], # Set the color using the predefined palette
    xlabel="Sample Index", # Label for the x-axis
    ylabel="Amplitude" # Label for the y-axis
)

# Display the plot
plt.show()
```



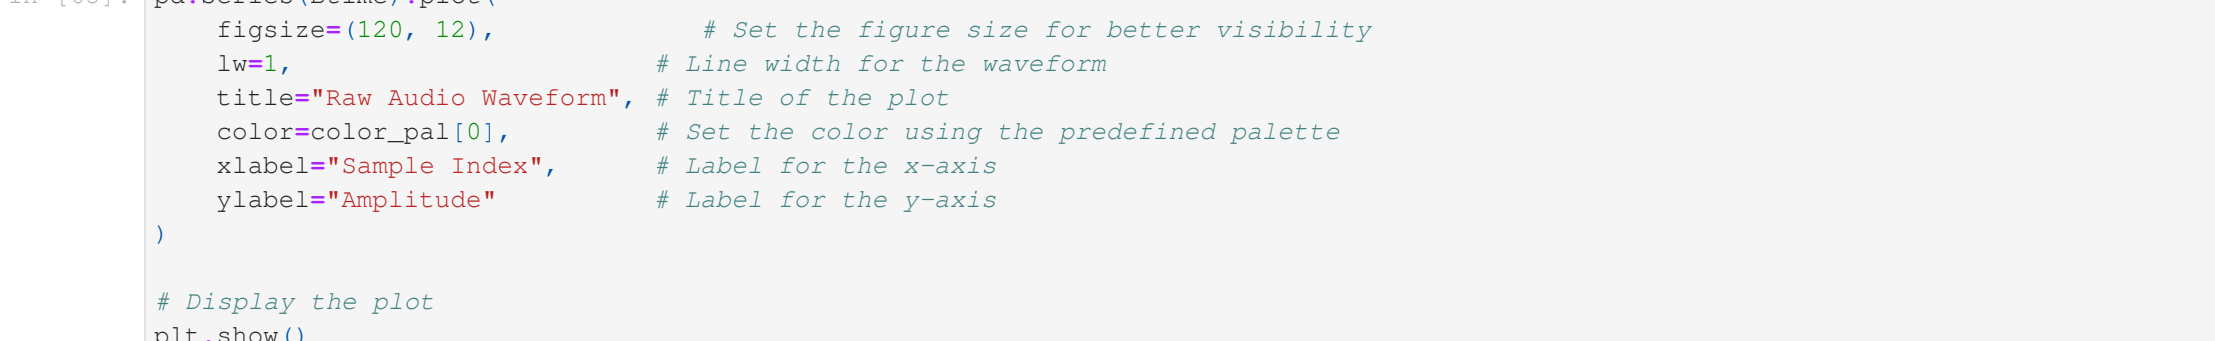
```
In [62]: #Plotting the raw audio waveform for the B note
pd.Series(Btime).plot(
    figsize=(120, 12), # Set the figure size for better visibility
    lw=1, # Line width for the waveform
    title="Raw Audio Waveform", # Title of the plot
    color=color_pal[0], # Set the color using the predefined palette
    xlabel="Sample Index", # Label for the x-axis
    ylabel="Amplitude" # Label for the y-axis
)

# Display the plot
plt.show()
```




```
In [63]: #for the B note
pd.Series(Btime).plot(
    figsize=(12, 7), # Set the figure size for better visibility
    lw=1, # Line width for the waveform
    title="Raw Audio Waveform", # Title of the plot
    color=color_pal[0], # Set the color using the predefined palette
    xlabel="Sample Index", # Label for the x-axis
    ylabel="Amplitude" # Label for the y-axis
)

# Display the plot
plt.show()
```



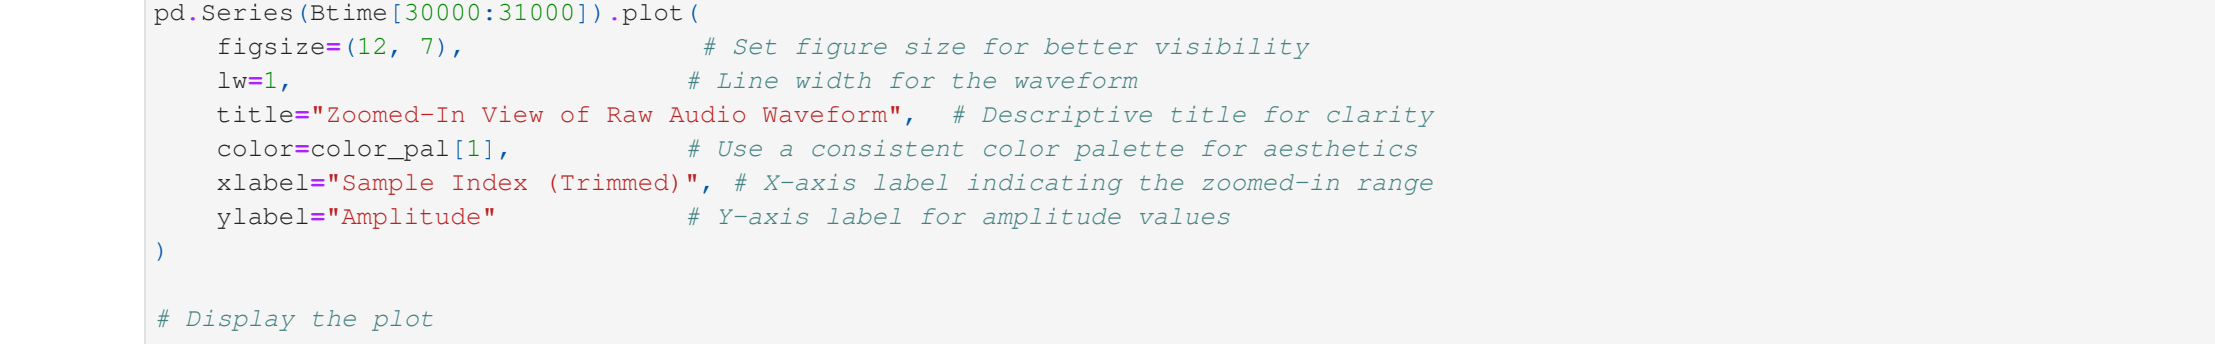
```
In [65]: #Plotting a trimmed section of the A note
pd.Series(Atime[30000:31000]).plot(
    figsize=(12, 7), # Set the figure size for better visibility
    lw=1, # Line width for the waveform
    title="Zoomed-In View of Raw Audio Waveform", # Descriptive title for clarity
    color=color_pal[0], # Use a consistent color palette for aesthetics
    xlabel="Sample Index (Trimmed)", # X-axis label indicating the zoomed-in range
    ylabel="Amplitude" # Y-axis label for amplitude values
)

# Display the plot
plt.show()
```



```
In [67]: #Plotting a trimmed section of the B note
pd.Series(Btime[30000:31000]).plot(
    figsize=(12, 7), # Set the figure size for better visibility
    lw=1, # Line width for the waveform
    title="Zoomed-In View of Raw Audio Waveform", # Descriptive title for clarity
    color=color_pal[0], # Use a consistent color palette for aesthetics
    xlabel="Sample Index (Trimmed)", # X-axis label indicating the zoomed-in range
    ylabel="Amplitude" # Y-axis label for amplitude values
)

# Display the plot
plt.show()
```



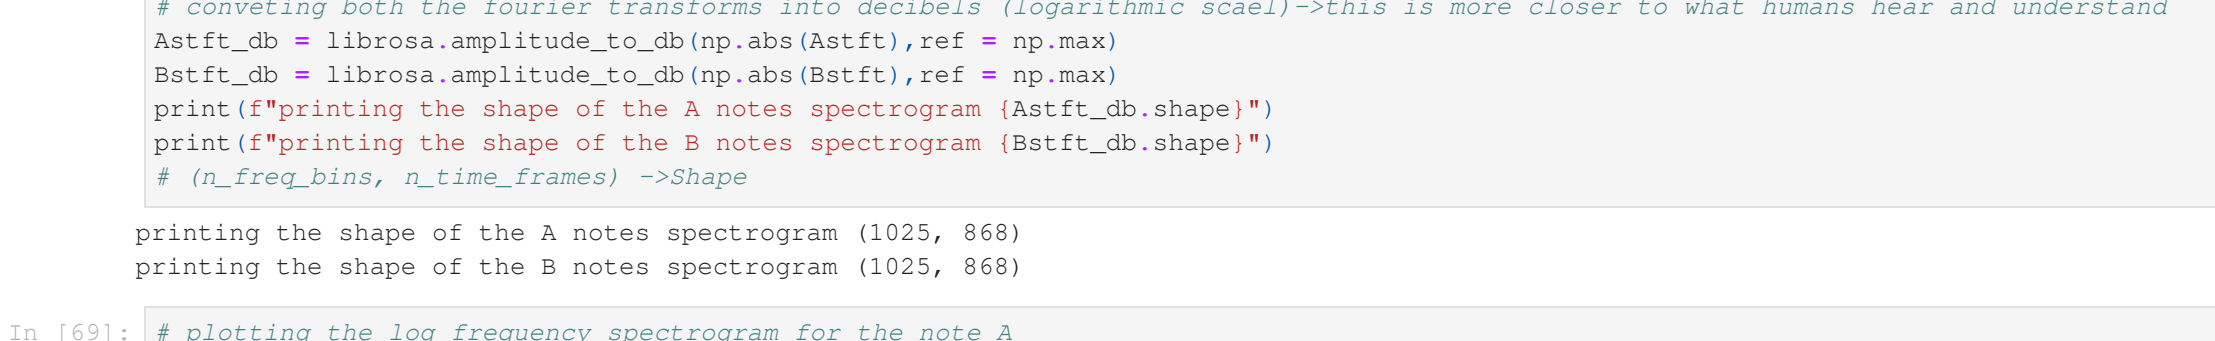
```
In [68]: #now taking the short time fourier transforms of both the notes
#this is to convert the graph that we have from a time domain to a frequency domain to then further use in a spectrograph
#converting both the fourier transforms into decibels (logarithmic scale)-this is more closer to what humans hear and understand
Bstft = librosa.stft(Btime)
Astft = librosa.stft(Atime)
Bstft_db = librosa.amplitude_to_db(np.abs(Astft),ref=np.max)
Astft_db = librosa.amplitude_to_db(np.abs(Bstft),ref=np.max)
print(f"Printing the shape of the note A in decibels",Bstft_db.shape)
print(f"Printing the shape of the B notes spectrogram (Bstft_db.shape)")
print(f"(n_freq_bins, n_time_frames) ->Shape
```

```
printing the shape of the A notes spectrogram (1025, 868)
printing the shape of the B notes spectrogram (1025, 868)
```

```
In [69]: # plotting the log frequency spectrogram for the note A
fig,ax = plt.subplots(figsize=(12,7))
img = librosa.display.specshow(Astft_db, x_axis="time",y_axis="log", sr = Asr,cmap = "magma", ax=ax)
ax.set_title("Log frequency spectrogram of the note A in decibels", fontsize=18, fontweight='bold')
ax.set_xlabel("Time (s)", fontsize=14)
ax.set_ylabel("Frequency (Hz)", fontsize=14)

# Add a color bar to indicate decibel levels
cbar = fig.colorbar(img, ax=ax, format="%2.0f dB")
cbar.set_label("Amplitude (dB)", fontsize=12)
```

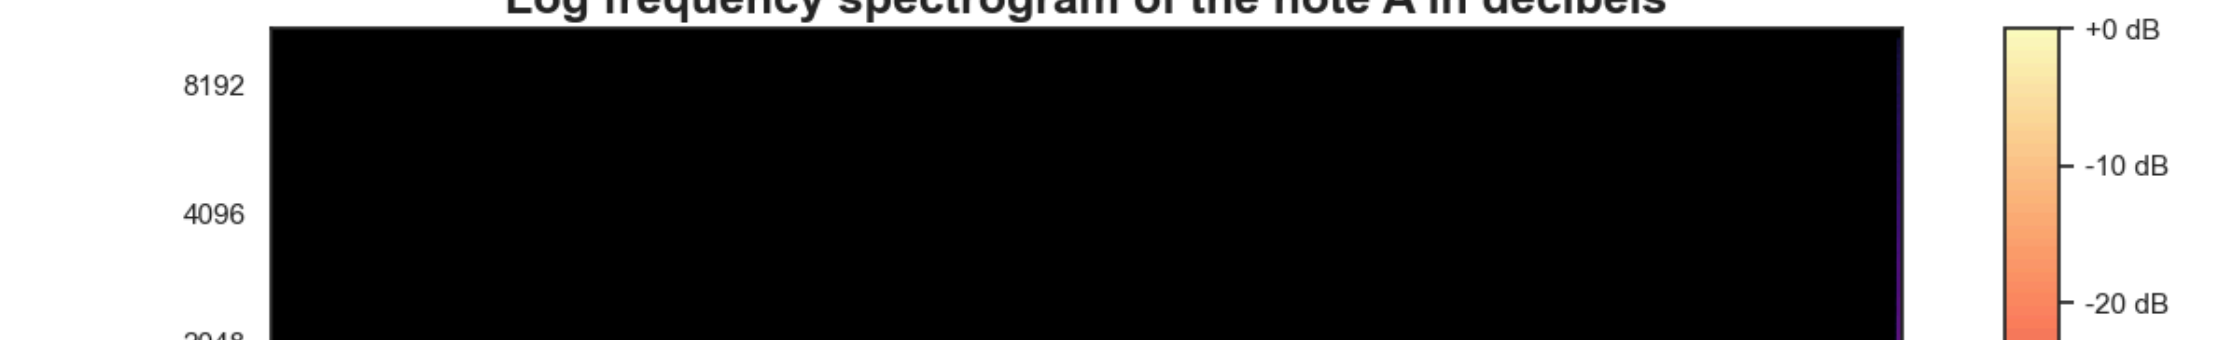
```
# Show the plot
plt.tight_layout() # Adjust layout to prevent clipping
plt.show()
```



```
In [70]: # plotting the log frequency spectrogram for the note B
fig,ax = plt.subplots(figsize=(12,7))
img = librosa.display.specshow(Bstft_db, x_axis="time",y_axis="log", sr = Asr,cmap = "magma", ax=ax)
ax.set_title("Log frequency spectrogram of the note B in decibels", fontsize=18, fontweight='bold')
ax.set_xlabel("Time (s)", fontsize=14)
ax.set_ylabel("Frequency (Hz)", fontsize=14)

# Add a color bar to indicate decibel levels
cbar = fig.colorbar(img, ax=ax, format="%2.0f dB")
cbar.set_label("Amplitude (dB)", fontsize=12)
```


```
# Show the plot
plt.tight_layout() # Adjust layout to prevent clipping
plt.show()
```



```
In [71]: #Plotting the log scale mel spectrogram for the note A
fig,ax = plt.subplots(figsize=(12,7))
img = librosa.display.specshow(A_db_mel, sr=Asr
    x_axis="time",
    y_axis="mel",
    cmap="magma",
    ax=ax)
ax.set_title("Mel Spectrogram for A (Log Scale in Decibels)", fontsize=18, fontweight="bold")
ax.set_xlabel("Time", fontsize=14)
ax.set_ylabel("Mel Frequency (Hz)", fontsize=14)

# Add a Color Bar to Show the Decibel Range
cbar = fig.colorbar(img, ax=ax, format="%2.0f dB")
cbar.set_label("Amplitude (dB)", fontsize=12)
```

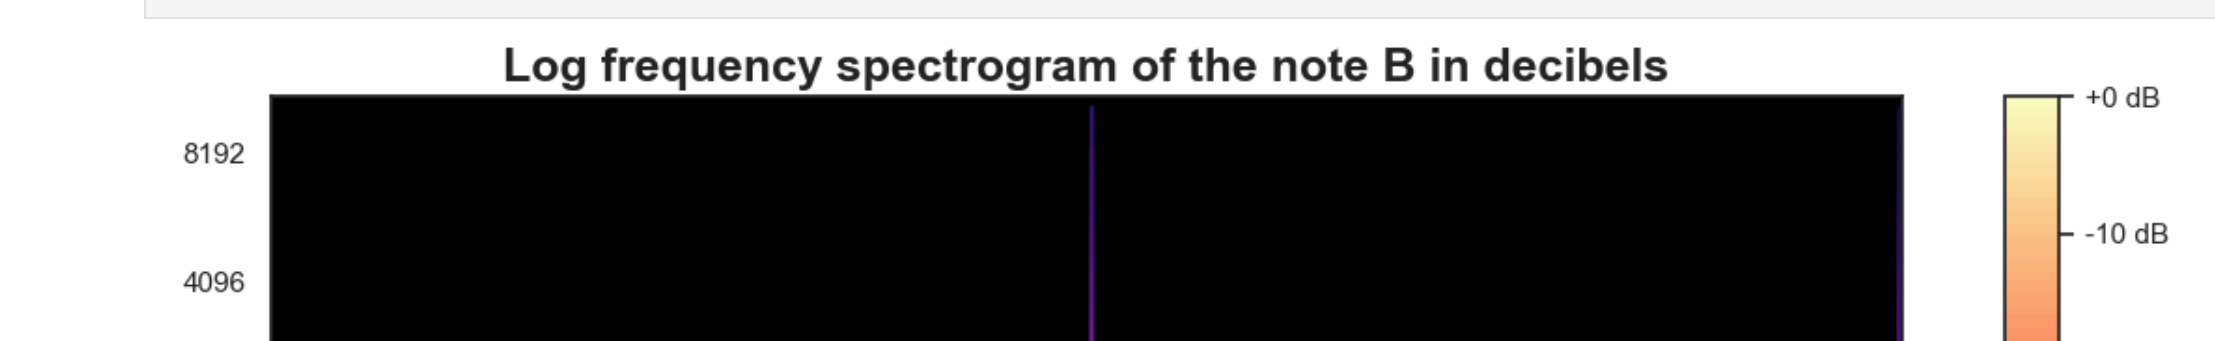
```
# Adjust Layout for Better Visualization
plt.tight_layout()
plt.show()
```



```
In [72]: #Plotting the log scale mel spectrogram for the note B
fig,ax = plt.subplots(figsize=(12,7))
img = librosa.display.specshow(B_db_mel, sr=Asr
    x_axis="time",
    y_axis="mel",
    cmap="magma",
    ax=ax)
ax.set_title("Mel Spectrogram for B (Log Scale in Decibels)", fontsize=18, fontweight="bold")
ax.set_xlabel("Time", fontsize=14)
ax.set_ylabel("Mel Frequency (Hz)", fontsize=14)

# Add a Color Bar to Show the Decibel Range
cbar = fig.colorbar(img, ax=ax, format="%2.0f dB")
cbar.set_label("Amplitude (dB)", fontsize=12)
```

```
# Adjust Layout for Better Visualization
plt.tight_layout()
plt.show()
```



```
In [73]: # Convert the Mel Spectrogram to Decibels (Logarithmic Scale)
A_db_mel = librosa.amplitude_to_db(A_mel_spectrogram_shape, ref=np.max)

# Display Information
print(f"Mel Spectrogram (in dB) Shape for A: {A_db_mel.shape}")
print(f"Decibel Range: {A_db_mel.min():.2f} dB to {A_db_mel.max():.2f} dB")

Mel Spectrogram (in dB) Shape for A: (128, 868)
Decibel Range: -80.00 dB to 0.00 dB
```

```
In [74]: # Convert the Mel Spectrogram to Decibels (Logarithmic Scale)
B_db_mel = librosa.amplitude_to_db(B_mel_spectrogram_shape, ref=np.max)

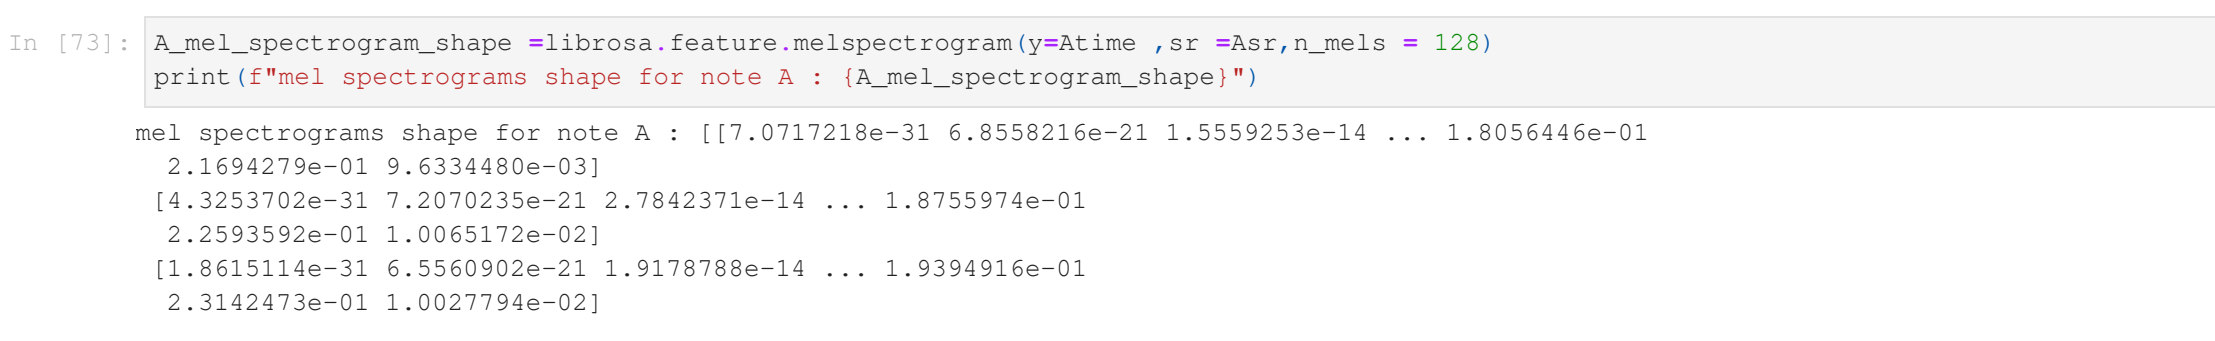
# Display Information
print(f"Mel Spectrogram (in dB) Shape: {B_db_mel.shape}")
print(f"Decibel Range: {B_db_mel.min():.2f} dB to {B_db_mel.max():.2f} dB")

Mel Spectrogram (in dB) Shape: (128, 868)
Decibel Range: -80.00 dB to 0.00 dB
```

```
In [75]: #Plotting the log scale mel spectrogram for the A note
fig,ax = plt.subplots(figsize=(12,7))
img = librosa.display.specshow(A_db_mel, sr=Asr
    x_axis="time",
    y_axis="mel",
    cmap="magma",
    ax=ax)
ax.set_title("Mel Spectrogram for A (Log Scale in Decibels)", fontsize=18, fontweight="bold")
ax.set_xlabel("Time", fontsize=14)
ax.set_ylabel("Mel Frequency (Hz)", fontsize=14)

# Add a Color Bar to Show the Decibel Range
cbar = fig.colorbar(img, ax=ax, format="%2.0f dB")
cbar.set_label("Amplitude (dB)", fontsize=12)
```

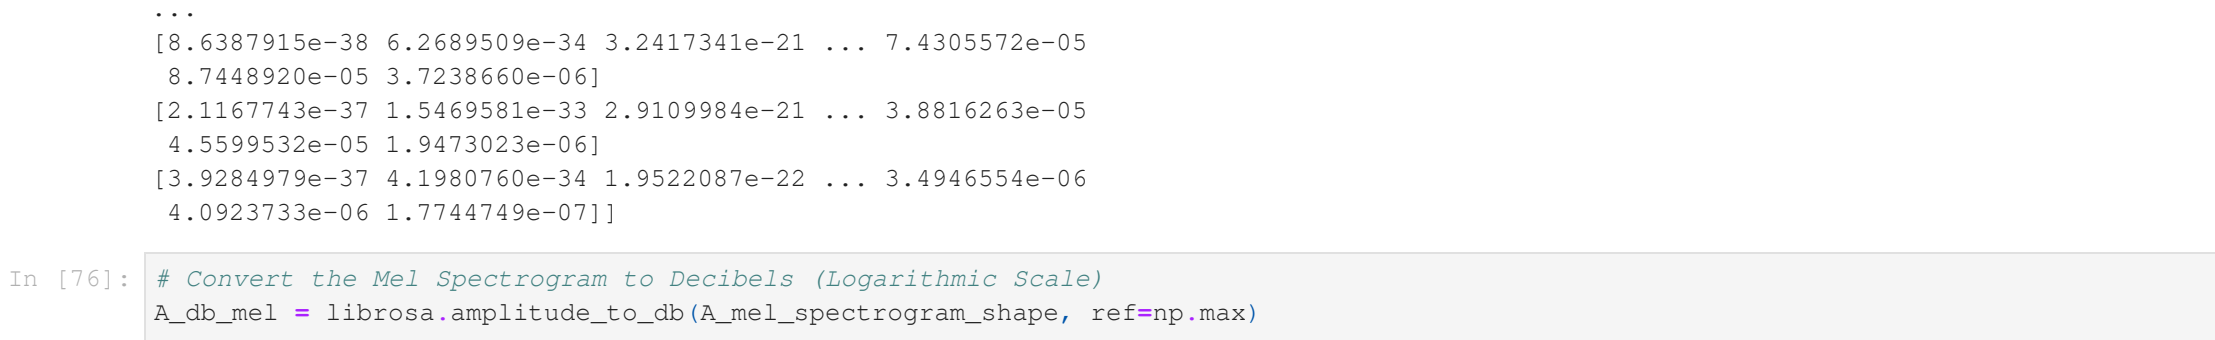
```
# Adjust Layout for Better Visualization
plt.tight_layout()
plt.show()
```



```
In [80]: #Plotting the log scale mel spectrogram for the B note
fig,ax = plt.subplots(figsize=(12,7))
img = librosa.display.specshow(B_db_mel, sr=Asr
    x_axis="time",
    y_axis="mel",
    cmap="magma",
    ax=ax)
ax.set_title("Mel Spectrogram for B (Log Scale in Decibels)", fontsize=18, fontweight="bold")
ax.set_xlabel("Time", fontsize=14)
ax.set_ylabel("Mel Frequency (Hz)", fontsize=14)

# Add a Color Bar to Show the Decibel Range
cbar = fig.colorbar(img, ax=ax, format="%2.0f dB")
cbar.set_label("Amplitude (dB)", fontsize=12)
```

```
# Adjust Layout for Better Visualization
plt.tight_layout()
plt.show()
```



```
In [81]: print("A note on decibels: The scale is a relative scale, 0db is the minimum the human ear can hear,-ve decibels indicate that the sound is very little in amplitude and not a negative sound(because ofc it isn't)
also in digital systems 0db is the loudest sound you can get before any distortion is induced into the audio
```

```
In [82]: print("The original plan was to get multiple of these files that carried note sounds and base my classification on that but i dont have the original plan was to get multiple of these files that carried note sounds and base my classification on that but i dont have that much of storage on my system
so i'm trying to create snippets of the two audio files and feed that to the dataset
```

```
In [104]: snippet_duration = 3 # seconds
snippets_A = []
snippets_B = []
for i in range(0, len(Atime), samples_per_snippet_A):
    snippet_A = Atime[i:i+samples_per_snippet_A]
    snippets_A.append(snippet_A)
for i in range(0, len(Btime), samples_per_snippet_B):
    snippet_B = Btime[i:i+samples_per_snippet_B]
    snippets_B.append(snippet_B)
```

```
# Extract MFCC features from each snippet
features_A = librosa.feature.mfcc(y=snippets_A, sr=Asr, n_mfcc=13).mean(axis=1)
features_B = librosa.feature.mfcc(y=snippets_B, sr=Br, n_mfcc=13).mean(axis=1)
```

```
# Combine the features and create corresponding labels
X = np.array(features_A + features_B)
y = np.array([0]*len(features_A) + [1]*len(features_B))
```

```
# Split the data so we can test the model's performance on unseen data.
X_train,X_test,y_train,y_test = train_test_split(X,y, test_size=0.3, random_state=42)
```

```
# It's important to scale the features so the model can train properly.
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
# Let's use a simple K-Nearest Neighbors (KNN) classifier.
model = KNeighborsClassifier(n_neighbors=3)
model.fit(X_train_scaled,y_train)
```

```
# Make predictions and check the model's accuracy.
y_pred = model.predict(X_test_scaled)
print(f"Model Classification Report")
print(f"Accuracy: {accuracy_score(y_test,y_pred):.2f}")
print(classification_report(y_test,y_pred, target_names=('Note A', 'Note B')))
```

```
# Let's try to classify a snippet from the original files.
new_snippet_path = audio_file1[0] # Or audio_file2[0]
new_snippet = librosa.load(new_snippet_path,duration=3) # Load just a 3-second snippet
```

```
# Process the new snippet in the same way as the training data.
new_snippet_features = librosa.feature.mfcc(y=new_snippet, sr=Asr, n_mfcc=13).mean(axis=1)
new_snippet_features = new_snippet_features.reshape(1,-1) # Reshape for the model
```

```
# Scale in the new snippet's features using the same scaler
new_snippet_scaled = scaler.transform(new_snippet_features)
```

```
# Make a prediction
prediction = model.predict(new_snippet_scaled)
predicted_class = "Note A" if prediction[0] == 0 else "Note B"
```

```
print(f"Prediction for the new snippet: {predicted_class}")
ipd.Audio(new_snippet,rate=Asr)
```

```
Model Evaluation:
Created 6 snippets for Note A.
Accuracy: 1.00
Classification Report:
```

	precision	recall	f1-score	support
Note A	1.00	1.00	1.00	1
Note B	1.00	1.00	1.00	3
accuracy avg	1.00	1.00	1.00	4
weighted avg	1.00	1.00	1.00	4

```
Prediction for the new snippet: Note A
```

```
Out[104]: 
```

```
In [106]: #checking with the second file
new_snippet_path = audio_file2[0] # Or audio_file1[0]
new_snippet, _ = librosa.load(new_snippet_path,duration=3) # Load just a 3-second snippet

# Process the new snippet in the same way as the training data.
new_snippet_features = librosa.feature.mfcc(y=new_snippet, sr=Asr, n_mfcc=13).mean(axis=1)
new_snippet_features = new_snippet_features.reshape(1,-1) # Reshape for the model
```

```
# Scale in the new snippet's features using the same scaler
new_snippet_scaled = scaler.transform(new_snippet_features)
```

```
# Make a prediction
prediction = model.predict(new_snippet_scaled)
predicted_class = "Note A" if prediction[0] == 0 else "Note B"
```

```
print(f"Prediction for the new snippet: {predicted_class}")
ipd.Audio(new_snippet,rate=Asr)
```

```
Prediction for the new snippet: Note B
```

```
Out[106]: 
```

```
In [111]:
```